



(a) All-CNN architecture



Each block: convolution $\rightarrow$ batch normalization $\rightarrow$ ReLU. Tensor labels: channels $\times$ spatial size. No fully connected layers.

(b) Training protocol

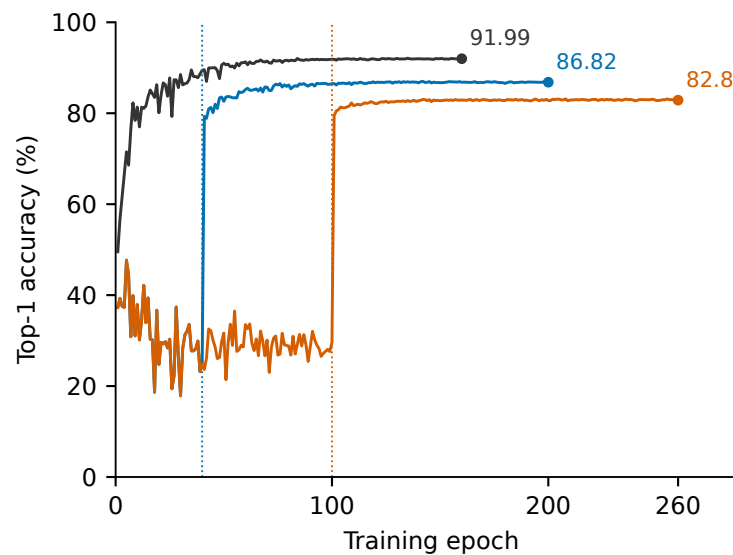


Hatched: bilinear $32 \rightarrow 8 \rightarrow 32$ blur

CIFAR-10 · independent reconstruction · training seed 0 only. Fisher: 300 clear training probes, mean of 5 posterior-label draws.

Final-model measurements, not Fisher trajectories during training. One seed does not establish a permanent deficit or a causal mechanism.

(c) Clear-test accuracy



(d) Final-model Fisher

